



Machine Learning 2D5362

Lecture 1: Introduction to Machine Learning



Machine Learning

- Date/Time:
Tuesday ???
Thursday 13.30
- Location:
BB2 ?
- Course requirements:
active participation
homework assignments
course project
- Credits:
3-5 credits depending on course project
- Course webpage:
<http://www.nada.kth.se/~hoffmann/ml.html>



Course Material

Textbook (recommended):

- Machine Learning

Tom M. Mitchell, McGraw Hill, 1997

ISBN: 0-07-042807-7 (available as paperback)

Further readings:

- An Introduction to Genetic Algorithms

Melanie Mitchell, MIT Press, 1996

- Reinforcement Learning – An Introduction

Richard Sutton, MIT Press, 1998

- Selected publications:

check course webpage



Course Overview

- Introduction to machine learning
- Concept learners
- Decision tree learning
- Neural networks
- Evolutionary algorithms
- Instance based learning
- Reinforcement learning
- Machine learning in robotics



Software Packages & Datasets

- MLC++
 - Machine learning library in C++
 - <http://www.sig.com/Technology/mlc>
- GALIB
 - MIT GALib in C++
 - <http://lancet.mit.edu/ga>
- UCI
 - Machine Learning Data Repository UC Irvine
 - <http://www.ics.uci.edu/~mlearn/ML/Repository.html>



Possible Course Projects

- Apply machine learning techniques to your own problem e.g. classification, clustering, data modeling, object recognition
- Investigating combining multiple classifiers
- Comparing different approaches in genetic fuzzy systems
- Learning robotic behaviors using evolutionary techniques or reinforcement learning
 - LEGO Mindstorm
 - Scout


Scout Robots

- 16 Sonar sensors
- Laser range scanner
- Odometry
- Differential drive
- Simulator
- API in C



LEGO Mindstorms

- Touch sensor
- Light sensor
- Rotation sensor
- Video cam
- Motors





Learning & Adaptation

- "Modification of a behavioral tendency by expertise."
(Webster 1984)
- "A learning machine, broadly defined is any device whose actions are influenced by past experiences." (Nilsson 1965)
- "Any change in a system that allows it to perform better the second time on repetition of the same task or on another task drawn from the same population." (Simon 1983)
- "An improvement in information processing ability that results from information processing activity." (Tanimoto 1990)



Learning

Definition:

A computer program is said to **learn** from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience.



Disciplines relevant to ML

- Artificial intelligence
- Bayesian methods
- Control theory
- Information theory
- Computational complexity theory
- Philosophy
- Psychology and neurobiology
- Statistics



Applications of ML


- Learning to recognize spoken words
 - SPHINX (Lee 1989)
- Learning to drive an autonomous vehicle
 - ALVINN (Pomerleau 1989)
- Learning to classify celestial objects
 - (Fayyad et al 1995)
- Learning to play world-class backgammon
 - TD-GAMMON (Tesauro 1992)
- Designing the morphology and control structure of electro-mechanical artefacts
 - GOLEM (Lipton, Pollock 2000)

Artificial Life


■ GOLEM Project (Nature: Lipson, Pollack 2000)

<http://golem03.cs-i.brandeis.edu/index.html>


- Evolve simple electromechanical locomotion machines from basic building blocks (bars, actuators, artificial neurons) in a simulation of the physical world (gravity, friction).
- The individuals that demonstrate the best locomotion ability are fabricated through rapid prototyping technology.



Evolvable Robot



Arrow



Ratchet



Tetra





Evolved Creatures

Evolved creatures: Sims (1994)

<http://genarts.com/karl/evolved-virtual-creatures.html>

Darwinian evolution of virtual block creatures for swimming, jumping, following, competing for a block

**Evolved Virtual
Creatures**

**Examples from
work in progress**



Learning Problem

Learning: improving with experience at some task

- Improve over task T
- With respect to performance measure P
- Based on experience E

Example: Learn to play checkers:

- T : play checkers
- P : percentage of games won in a tournament
- E : opportunity to play against itself



Learning to play checkers

- T: play checkers
- P: percentage of games won
- What experience?
- What exactly should be learned?
- How shall it be represented?
- What specific algorithm to learn it?



Type of Training Experience

- Direct or indirect?
 - Direct: board state -> correct move
 - Indirect: outcome of a complete game
 - Credit assignment problem
- Teacher or not ?
 - Teacher selects board states
 - Learner can select board states
- Is training experience representative of performance goal?
 - Training playing against itself
 - Performance evaluated playing against world champion



Choose Target Function

- **ChooseMove** : $B \rightarrow M$: board state \rightarrow move
 - Maps a legal board state to a legal move
- **Evaluate** : $B \rightarrow V$: board state \rightarrow board value
 - Assigns a numerical score to any given board state, such that better board states obtain a higher score
 - Select the best move by evaluating all successor states of legal moves and pick the one with the maximal score



Possible Definition of Target Function

- If b is a final board state that **is won** then $V(b) = 100$
- If b is a final board state that **is lost** then $V(b) = -100$
- If b is a final board state that **is drawn** then $V(b)=0$
- If b is not a final board state, then $V(b)=V(b')$, where b' is the best final board state that can be achieved starting from b and playing optimally until the end of the game.
- Gives correct values but is not operational

State Space Search




eu sou
vermelho

jogador vermelho a jogar


$$V(b) = ?$$

$$V(b) = \max_i V(b_i)$$

$m_1 : b \rightarrow b_1$ $m_2 : b \rightarrow b_2$ $m_3 : b \rightarrow b_3$



State Space Search









$V(b_1) = ?$



$V(b_1) = \min_i V(b_i)$

$m_4 : b \rightarrow b_4$ $m_5 : b \rightarrow b_5$ ↓



$m_6 : b \rightarrow b_6$





Final Board States



Black wins: $V(b)=-100$




Red wins: $V(b)=100$



draw: $V(b)=0$








Depth-First Search






Breadth-First Search



Number of Board States

Tic-Tac-Toe:


$$\begin{aligned} \# \text{board states} &< 9!/(5! 4!) + 9!/(1! 4! 4!) + \dots \\ &\dots + 9!/(2! 4! 3!) + \dots 9 = 6045 \end{aligned}$$



4 x 4 checkers: (no queens)

#board states = ?

$$\# \text{board states} < 8 \times 7 \times 6 \times 5 \times 2^2 / (2! \times 2!) = 1680$$



Regular checkers (8x8 board, 8 pieces each)

$$\# \text{board states} < 32! \times 2^{16} / (8! \times 8! \times 16!) = 5.07 \times 10^{17}$$



Choose Representation of Target Function

- Table look-up
- Collection of rules
- Neural networks
- Polynomial function of board features
- Trade-off in choosing an expressive representation:
 - Approximation accuracy
 - Number of training examples to learn the target function

Representation of Target Function

$$V(b) = \omega_0 + \omega_1 bp(b) + \omega_2 rp(b) + \omega_3 bk(b) + \omega_4 rk(b) + \omega_5 bt(b) + \omega_6 rt(b)$$

- $bp(b)$: #black pieces
- $rp(b)$: #red pieces
- $bk(b)$: #black kings
- $rk(b)$: #red kings
- $bt(b)$: #red pieces threatened by black
- $rt(b)$: #black pieces threatened by red



Obtaining Training Examples

- $V(b)$: true target function
- $V'(b)$: learned target function
- $V_{\text{train}}(b)$: training value
- Rule for estimating training values:
- $V_{\text{train}}(b) \leftarrow V'(\text{Successor}(b))$

$\begin{pmatrix} -100 \\ 100 \\ 0 \end{pmatrix}$



Choose Weight Training Rule

LMS weight update rule:

- Select a training example b at random

1. Compute error(b)

actual

anterior

$$\text{error}(b) = V_{\text{train}}(b) - V'(b)$$

2. For each board feature f_i , update


$$\text{weight } \omega_i \leftarrow \omega_i + \eta f_i \text{error}(b)$$

η : learning rate approx. 0.1

Example: 4x4 checkers

$$V(b) = \omega_0 + \omega_1 \text{rp}(b) + \omega_2 \text{bp}(b)$$


Initial weights: $\omega_0 = -10$, $\omega_1 = 75$, $\omega_2 = -60$



$$V(b_0) = \omega_0 + \omega_1 * 2 + \omega_2 * 2 = 20$$


$\mathbf{m}_1 : b \rightarrow b_1$

$$V(b_1) = 20$$




$\mathbf{m}_2 : b \rightarrow b_2$

$$V(b_2) = 20$$




$\mathbf{m}_3 : b \rightarrow b_3$


$$V(b_3) = 20$$



Example 4x4 checkers



$V(b_0)=20$



$V(b_1)=20$

1. Compute $\text{error}(b_0) = V_{\text{train}}(b) - V(b_0) = V(b_1) - V(b_0) = 0$
2. For each board feature f_i , update weight


$$\omega_i \leftarrow \omega_i + \eta f_i \text{error}(b)$$

$$\omega_0 \leftarrow \omega_0 + 0.1 * 1 * 0$$


$$\omega_1 \leftarrow \omega_1 + 0.1 * 2 * 0$$

$$\omega_2 \leftarrow \omega_2 + 0.1 * 2 * 0$$


Example: 4x4 checkers




→
 $V(b_0)=20$



$V(b_1)=20$




←
 $V(b_3)=20$




$V(b_2)=20$


Example: 4x4 checkers



$V(b_3)=20$




$V(b_{4a})=20$




$V(b_{4b})=-55$

Example 4x4 checkers



$V(b_3)=20$



$V(b_4)=-55$

1. Compute $\text{error}(b_3) = V_{\text{train}}(b) - V(b_3) = V(b_4) - V(b_3) = -75$
2. For each board feature f_i , update weight

$$\omega_i \leftarrow \omega_i + \eta f_i \text{error}(b) : \omega_0 = -10, \omega_1 = 75, \omega_2 = -60$$

$$\omega_0 \leftarrow \omega_0 - 0.1 * 1 * \cancel{75}, \omega_0 = -17.5$$


$$\omega_1 \leftarrow \omega_1 - 0.1 * 2 * 75, \omega_1 = 60$$

$$\omega_2 \leftarrow \omega_2 - 0.1 * 2 * 75, \omega_2 = -75$$


$f_i == 1$

Example: 4x4 checkers

$$\omega_0 = -17.5, \omega_1 = 60, \omega_2 = -75$$




$$V(b_4) = -107.5$$




$$V(b_5) = -107.5$$

Example 4x4 checkers



$V(b_5) = -107.5$



$V(b_6) = -167.5$

$$\text{error}(b_5) = V_{\text{train}}(b) - V(b_5) = V(b_6) - V(b_5) = -60$$

$$\omega_0 = -17.5, \omega_1 = 60, \omega_2 = -75$$


$$\omega_i \leftarrow \omega_i + \eta f_i \text{error}(b)$$

$$\omega_0 \leftarrow \omega_0 - 0.1 * 1 * 60, \omega_0 = -23.5$$

$$\omega_1 \leftarrow \omega_1 - 0.1 * 1 * 60, \omega_1 = 54$$

$$\omega_2 \leftarrow \omega_2 - 0.1 * 2 * 60, \omega_2 = -87$$

Example 4x4 checkers



Final board state: black won $V_f(b) = -100$

$$V(b_6) = -197.5$$

$$\text{error}(b_6) = V_{\text{train}}(b) - V(b_6) = V_f(b_6) - V(b_6) = 97.5$$

$$\omega_0 = -23.5, \omega_1 = 54, \omega_2 = -87$$

$$\omega_i \leftarrow \omega_i + \eta f_i \text{error}(b)$$

$$\omega_0 \leftarrow \omega_0 + 0.1 * 1 * 97.5, \omega_0 = -13.75$$


$$\omega_1 \leftarrow \omega_1 + 0.1 * 0 * 97.5, \omega_1 = 54$$

$$\omega_2 \leftarrow \omega_2 + 0.1 * 2 * 97.5, \omega_2 = -67.5$$


Evolution of Value Function

Training data:
before ●
after ●


$$\omega = (-10, 75, -60)$$




$$\omega = (-17.5, 60, -75)$$



$$\omega = (-23.5, 54, -87)$$



$$\omega = (-13.75, 54, -67.5)$$





Design Choices

Determine Type of Training Experience

Games against experts

Games against self

Table of correct moves

Determine Target Function

Board \rightarrow Move

Board \rightarrow Value

Determine Representation of Learned Function

polynomial

Linear function of six features

Artificial neural network

Determine Learning Algorithm

Gradient descent

Linear programming



Learning Problem Examples

- Credit card applications
 - Task T: Distinguish "good" applicants from "risky" applicants.
 - Performance measure P : ?
 - Experience E : ? (direct/indirect)
 - Target function : ?



Performance Measure P:

- Error based: minimize percentage of incorrectly classified customers : $P = N_{fp} + N_{fn} / N$
 - N_{fp} : # false positives (rejected good customers)
 - N_{fn} : # false negatives (accepted bad customers)
- Utility based: maximize expected profit of credit card business: $P = N_{cp} * U_{cp} + N_{fn} * U_{fn}$
 - U_{cp} : expected utility of an accepted good customer
 - U_{fn} : expected utility/loss of an accepted bad customer




Experience E:

- Direct: Decisions on credit card applications made by a human financial expert
 - Training data: <customer inf., reject/accept>
- Direct: Actual customer behavior based on previously accepted customers
 - Training data: <customer inf., good/bad>
 - Problem: Distribution of applicants $P_{\text{applicant}}$ is not identical with training data P_{train}
- Indirect: Evaluate a decision policy based on the profit you made over the past N years.

Distribution of Applicants

Good customers


Bad customers




Distribution of Accepted Customers

Good customers

Bad customers



What is the optimal decision boundary?





Target Function

Customer record:

income, owns house, credit history, age, employed, accept
\$40000, yes, good, 38, full-time, yes
\$25000, no, excellent, 25, part-time, no
\$50000, no, poor, 55, unemployed, no

- T: Customer data → accept/reject
- T: Customer data → probability good customer
- T: Customer data → expected utility/profit



Learning methods

- Decision rules:
 - If income < \$30.000 then reject
- Bayesian network:
 - $P(\text{good} \mid \text{income, credit history,....})$
- Neural Network:
- Nearest Neighbor:
 - Take the same decision as for the customer in the data base that is most similar to the applicant



Learning Problem Examples

- Obstacle Avoidance Behavior of a Mobile Robot
 - Task T: Navigate robot safely through an environment.
 - Performance measure P : ?
 - Experience E : ?
 - Target function : ?



Performance Measure P:

- P: Maximize time until collision with obstacle
- P: Maximize distance travelled until collision with obstacle
- P: Minimize rotational velocity, maximize translational velocity
- P: Minimize error between control action of a human operator and robot controller in the same situation



Training Experience E:

- Direct: Monitor human operator and use her control actions as training data:
 - $E = \{ \langle \text{perception}_i, \text{action}_i \rangle \}$
- Indirect: Operate robot in the real world or in a simulation. Reward desirable states, penalize undesirable states
 - $V(b) = +1$ if $v > 0.5 \text{ m/s}$
 - $V(b) = +2$ if $\omega < 10 \text{ deg/s}$
 - $V(b) = -100$ if bumper state = 1

Question: Internal or external reward ?



Target Function

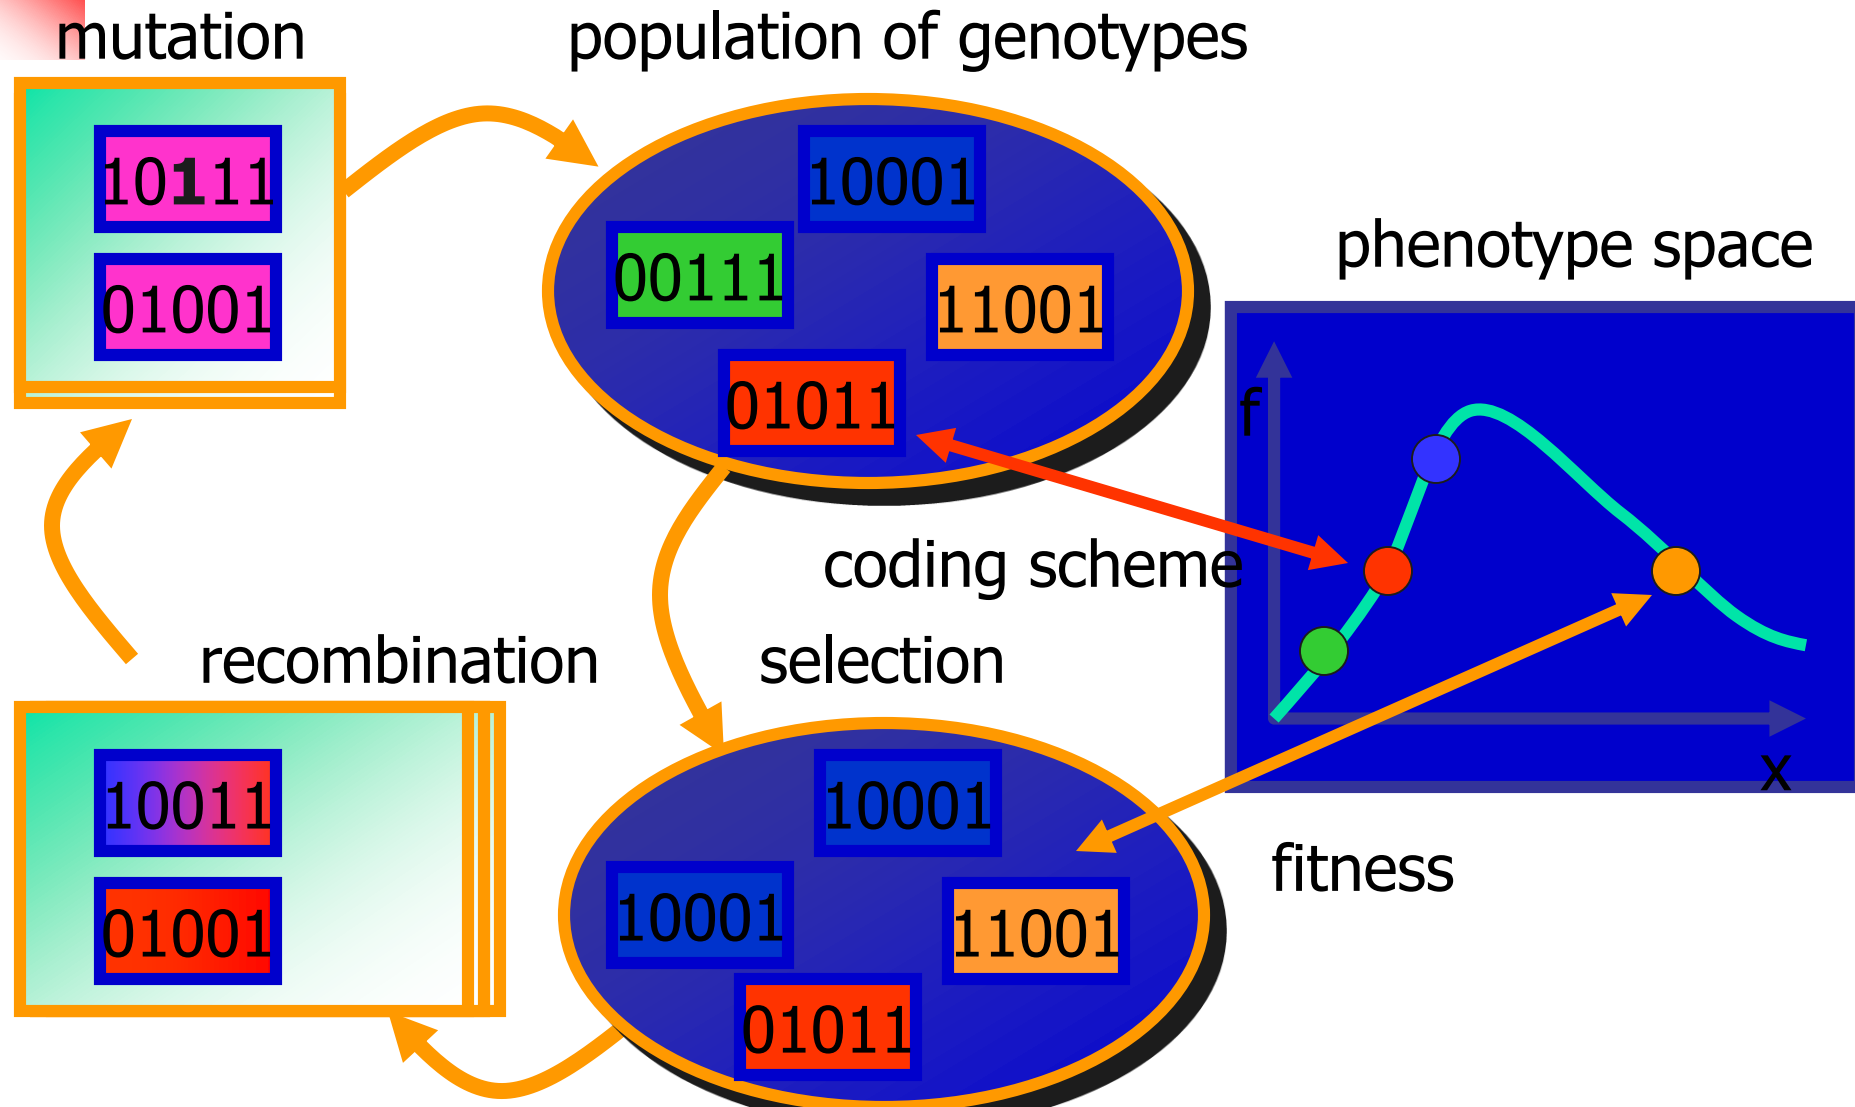
- Choose action:
 - A: perception \rightarrow action
 - Sonar readings: $s_1(t) \dots s_n(t) \rightarrow \langle v, \omega \rangle$
- Evaluate perception/state:
 - V: $s_1(t) \dots s_n(t) \rightarrow V(s_1(t) \dots s_n(t))$
 - Problem: states are only partially observable therefore world seems non-deterministic
 - Markov Decision Process : successor state $s(t+1)$ is a probabilistic function of current state $s(t)$ and action $a(t)$
- Evaluate state/action pairs:
 - V: $s_1(t) \dots s_n(t), a(t) \rightarrow V(s_1(t) \dots s_n(t), a(t))$



Learning Methods

- Neural Networks
 - Require direct training experience
- Reinforcement Learning
 - Indirect training experience
- Evolutionary Algorithms
 - Indirect training experience

Evolutionary Algorithms



Evolution of Simple Navigation





Issues in Machine Learning

- What algorithms can approximate functions well and when?
- How does the number of training examples influence accuracy?
- How does the complexity of hypothesis representation impact it?
- How does noisy data influence accuracy?
- What are the theoretical limits of learnability?



Machine vs. Robot Learning

Machine Learning

- Learning in vacuum
- Statistically well-behaved data
- Mostly off-line
- Informative feed-back
- Computational time not an issue
- Hardware does not matter
- Convergence proof

Robot Learning

- Embedded learning
- Data distribution not homogeneous
- Mostly on-line
- Qualitative and sparse feed-back
- Time is crucial
- Hardware is a priority
- Empirical proof



Learning in Robotics

- behavioral adaptation:
adjust the parameters of individual behaviors according to some direct feedback signal (e.g. adaptive control)
- evolutionary adaptation:
application of artificial evolution to robotic systems
- sensor adaptation:
adopt the perceptual system to the environment
(e.g. classification of different contexts, recognition)
- learning complex, deliberative behaviors:
unsupervised learning based on sparse feedback from the environment, credit assignment problem
(e.g. reinforcement learning)